

Generative Adversarial Networks: A Comprehensive Analysis of Architectures, Training Dynamics, and Advanced Variants

Based on Reference: <https://github.com/wiseodd/generative-models>

September 4, 2025

Abstract

This report provides an in-depth exploration of Generative Adversarial Networks (GANs), encompassing their foundational architecture, game-theoretic training dynamics, challenges, and solutions. It includes detailed comparisons of core GAN variants and elaborates on advanced models such as StyleGAN, CycleGAN, and BigGAN. Drawing from the referenced GitHub repository, which implements various GANs in PyTorch and TensorFlow, this document integrates equations, professional diagrams, comparative tables, and practical use cases to offer a thorough, professional analysis suitable for researchers and practitioners.

Contents

1	Introduction	2
2	GAN Architecture and Training Dynamics: Game-Theoretic Interpretation	2
2.1	Core Architecture	2
2.2	Training Dynamics and Game-Theoretic Interpretation	3
3	Nash Equilibrium in GANs, Training Issues, and Game-Theoretic Solutions	3
3.1	Nash Equilibrium in GANs	3
3.2	Issues Associated with GAN Training	3
3.3	Possible Solutions from a Game-Theory Perspective	4
4	Discussion and Comparison of Vanilla GAN, Conditional GAN, Coupled GAN, LAP-GAN, and Adversarially Learned Inference	4
4.1	Vanilla GAN	4
4.2	Conditional GAN (CGAN)	5
4.3	Coupled GAN (CoGAN)	5
4.4	LAPGAN (Laplacian Pyramid GAN)	5
4.5	Adversarially Learned Inference (ALI)	6
4.6	Comparative Analysis	6
5	Discussion of Advanced GAN Types: StyleGAN, CycleGAN, BigGAN	6
5.1	StyleGAN	6

5.2	CycleGAN	7
5.3	BigGAN	8
6	Conclusion	8

1 Introduction

Generative Adversarial Networks (GANs), introduced by Ian Goodfellow et al. in 2014, represent a paradigm shift in generative modeling. Unlike traditional models that explicitly estimate data distributions, GANs employ an adversarial framework where two neural networks compete: a generator creates synthetic data, and a discriminator evaluates its authenticity. This report elaborates on key aspects of GANs, leveraging implementations from the referenced repository (<https://github.com/wiseodd/generative-models>), which includes code for Vanilla GAN, Conditional GAN, and others. Generated samples are stored in directories like `GAN/{model}/out`. The discussion extends to advanced variants not directly in the repository but built upon similar principles.

2 GAN Architecture and Training Dynamics: Game-Theoretic Interpretation

2.1 Core Architecture

The GAN architecture comprises two primary components:

- **Generator (G):** A neural network that transforms random noise z , sampled from a prior distribution p_z (typically a multivariate Gaussian $\mathcal{N}(0, I)$), into a synthetic sample $G(z)$ that aims to mimic the real data distribution p_{data} . G is parameterized by θ_G and optimized to produce realistic outputs.
- **Discriminator (D):** A binary classifier that takes an input x (either real or generated) and outputs a scalar probability $D(x) \in [0, 1]$, indicating the likelihood that x is real. D is parameterized by θ_D and trained to maximize its accuracy in distinguishing real from fake samples.

This setup forms an adversarial loop, where G improves to deceive D, and D enhances to detect fakes. The architecture is illustrated below:

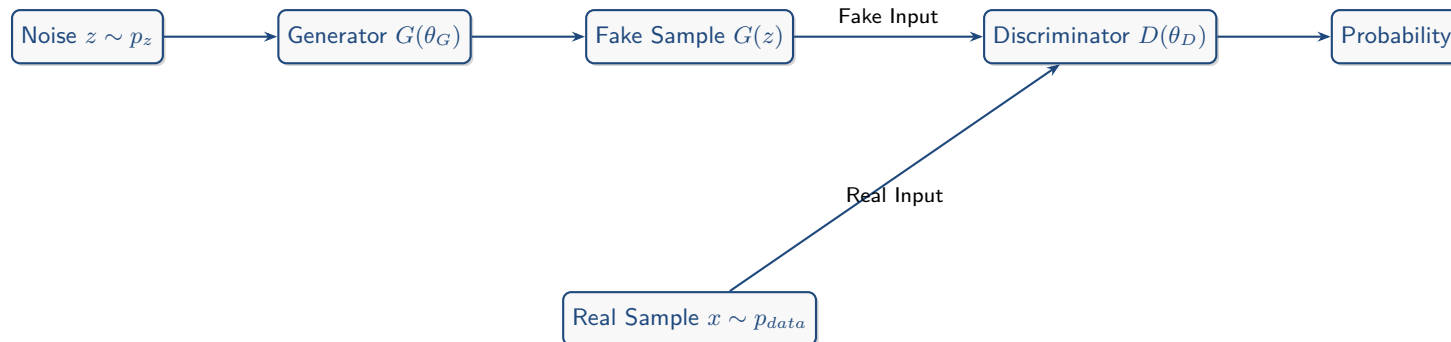


Figure 1: Vanilla GAN Architecture

In practice, G and D are deep neural networks, such as multi-layer perceptrons or convolutional networks for image data. The repository implements this in PyTorch and TensorFlow, with training loops alternating between D and G updates.

2.2 Training Dynamics and Game-Theoretic Interpretation

GAN training is conceptualized as a two-player, zero-sum minimax game in game theory. The value function is:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))] \quad (1)$$

From a game-theoretic perspective:

- This is a non-cooperative game where G and D are players with opposing objectives.
- The solution is a saddle point where neither player benefits from deviating unilaterally.
- Training involves stochastic gradient descent (SGD): Update θ_D via ascent on V , then θ_G via descent.
- Dynamics: Initially, D dominates if G produces poor samples, providing strong gradients for G. As G improves, equilibrium approaches, but instability can arise (detailed in Section 3).

Theoretically, under infinite capacity and data, G recovers p_{data} . Practically, training uses mini-batches and alternates k D steps per G step (often $k = 1$). The repository’s Vanilla GAN code demonstrates this loop, logging losses and saving samples.

3 Nash Equilibrium in GANs, Training Issues, and Game-Theoretic Solutions

3.1 Nash Equilibrium in GANs

In game theory, a Nash Equilibrium is a strategy profile where no player improves by changing strategy alone. In GANs, it occurs when:

- G produces samples from $p_g = p_{data}$.
- D is optimal, outputting $D^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_g(x)} = 0.5$ for all x (distributions indistinguishable).

At equilibrium, the value function relates to Jensen-Shannon Divergence (JSD):

$$V(G, D^*) = 2 \cdot \text{JSD}(p_{data} \| p_g) - \log 4 \quad (2)$$

Minimizing JSD to zero achieves equilibrium. This is a local Nash, but global convergence assumes convex-concave V , which neural networks violate, leading to challenges.

3.2 Issues Associated with GAN Training

GAN training is prone to several issues:

- **Mode Collapse:** G maps different z to the same output, capturing few modes of p_{data} while ignoring others. This arises from G exploiting D’s weaknesses.
- **Vanishing or Exploding Gradients:** If D is too strong, $D(G(z)) \approx 1$, making $\log(1 - D(G(z))) \approx -\infty$ (exploding); if weak, ≈ 0 (vanishing), stalling G.
- **Oscillations and Non-Convergence:** Losses fluctuate without settling, due to non-stationary objectives and saddle point navigation.
- **Hyperparameter Sensitivity:** Learning rates, batch sizes, and architectures critically affect stability; imbalance leads to divergence.
- **Evaluation Difficulty:** No single metric (e.g., Inception Score, FID) fully captures quality.

These stem from the game’s non-convex nature and finite model capacity.

3.3 Possible Solutions from a Game-Theory Perspective

Game theory offers strategies to mitigate issues:

- **Alternative Divergences (e.g., Wasserstein GAN - WGAN):** Replace JSD with Earth-Mover’s (Wasserstein) distance for smoother gradients:

$$\min_G \max_{D \in \mathcal{L}} \mathbb{E}_{x \sim p_{data}} [D(x)] - \mathbb{E}_{z \sim p_z} [D(G(z))] \quad (3)$$

where \mathcal{L} enforces 1-Lipschitz via clipping or penalties. Reduces mode collapse by providing gradients even when distributions don’t overlap.

- **Mode Regularization (e.g., Mode Regularized GAN - MRGAN):** Introduces a regularizer to encourage diversity, viewing collapse as suboptimal equilibria; penalizes low-variance generations.
- **Energy-Based Formulations (e.g., Energy-Based GAN - EBGAN):** Treats D as an energy function assigning low energy to real data, high to fakes, with margin losses for stability.
- **Multi-Agent Perspectives (e.g., Multiple Discriminators):** Use ensembles to approximate mixed strategies, avoiding pure strategy collapses.
- **Gradient Penalties and Improved Optimization:** Enforce constraints to approximate Nash, e.g., in Improved WGAN (WGAN-GP), adding:

$$\lambda \mathbb{E}_{\hat{x} \sim p_{\hat{x}}} [(\|\nabla_{\hat{x}} D(\hat{x})\|_2 - 1)^2] \quad (4)$$

where \hat{x} interpolates real and fake.

The repository implements WGAN, MRGAN, EBGAN, and others, providing code to experiment with these solutions.

4 Discussion and Comparison of Vanilla GAN, Conditional GAN, Coupled GAN, LAPGAN, and Adversarially Learned Inference

This section details the workings of selected GAN variants, followed by a comparative analysis.

4.1 Vanilla GAN

The foundational model (Equation 1). G generates unconditionally; D classifies globally. Suitable for simple distributions but unstable for complex data like images.

4.2 Conditional GAN (CGAN)

Extends Vanilla by conditioning on auxiliary information y (e.g., class labels, text):

$$\min_G \max_D V(D, G) = \mathbb{E}_{x, y \sim p_{data}} [\log D(x|y)] + \mathbb{E}_{z \sim p_z, y \sim p_y} [\log(1 - D(G(z|y)|y))] \quad (5)$$

- G inputs concatenated z and y ; D conditions on y for classification.
- Enhances control, e.g., generating specific digits in MNIST.

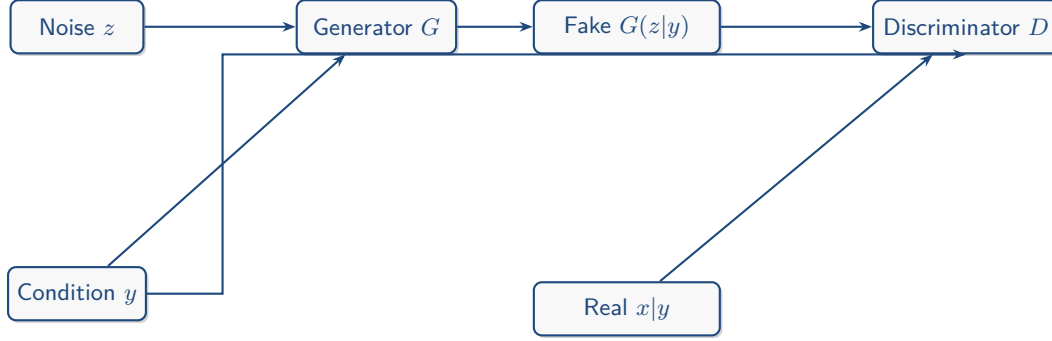


Figure 2: Conditional GAN Architecture

4.3 Coupled GAN (CoGAN)

Designed for learning joint distributions across related domains (e.g., color and grayscale images) without paired data. Two GANs (GAN1 for domain A, GAN2 for B) share weights in lower generator/discriminator layers to capture shared high-level features, diverging in higher layers for domain-specific details.

- Objective: Two independent minimax games, coupled via shared parameters.
- Training: Alternates updates, enforcing shared latent representations.

This promotes cross-domain consistency, useful for unpaired translation.

4.4 LAPGAN (Laplacian Pyramid GAN)

A hierarchical approach using Laplacian pyramids for progressive generation. Decomposes images into multi-scale residuals:

- At each pyramid level k , a conditional GAN upsamples low-res images and adds details.
- Objective: Per-level CGAN losses, conditioned on coarser images.
- Process: Start from coarse noise, refine iteratively to high-res.



Figure 3: LAPGAN Hierarchical Architecture

Reduces computational burden for high-resolution generation.

4.5 Adversarially Learned Inference (ALI)

Combines generation with inference by introducing an encoder E (or $q(z|x)$) to map data to latents. Discriminator judges joint pairs (x, z) :

$$\min_{G,E} \max_D V = \mathbb{E}_{x \sim p_{data}} [\log D(x, E(x))] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z), z))] \quad (6)$$

- Bidirectional: G maps $z \rightarrow x$, E maps $x \rightarrow z$.
- Improves mode coverage and enables posterior inference, unlike unidirectional GANs.

Repository includes implementations for Vanilla, CGAN, CoGAN, LAPGAN (as progressive), and ALI.

4.6 Comparative Analysis

Variant	Key Feature	Objective	Strengths	Limitations
Vanilla GAN	Unconditional generation	Eq. 1	Simple, foundational	Unstable, mode collapse
CGAN	Conditioned on labels/data	Eq. 5	Controllable outputs	Requires labeled data
CoGAN	Coupled multi-domain	Dual minimax	Unpaired domain learning	Limited to related domains
LAPGAN	Hierarchical pyramid	Per-level CGAN	High-res efficiency	Complex implementation
ALI	Bidirectional inference	Eq. 6	Inference capability	Increased computation

Table 1: Comparison of GAN Variants

Vanilla serves as baseline; CGAN adds control; CoGAN extends to multi-domain; LAPGAN scales resolution; ALI adds invertibility.

5 Discussion of Advanced GAN Types: StyleGAN, CycleGAN, BigGAN

This section elaborates on advanced GANs, including architectures, objectives, diagrams, and use cases.

5.1 StyleGAN

StyleGAN (Karras et al., 2019) revolutionizes image synthesis by disentangling latent factors for fine-grained control.

- **Architecture:** A mapping network (8-layer MLP) transforms noise z to intermediate latent w ; synthesis network uses w via AdaIN (Adaptive Instance Normalization) at each resolution level. Adds noise per layer for stochastic details.

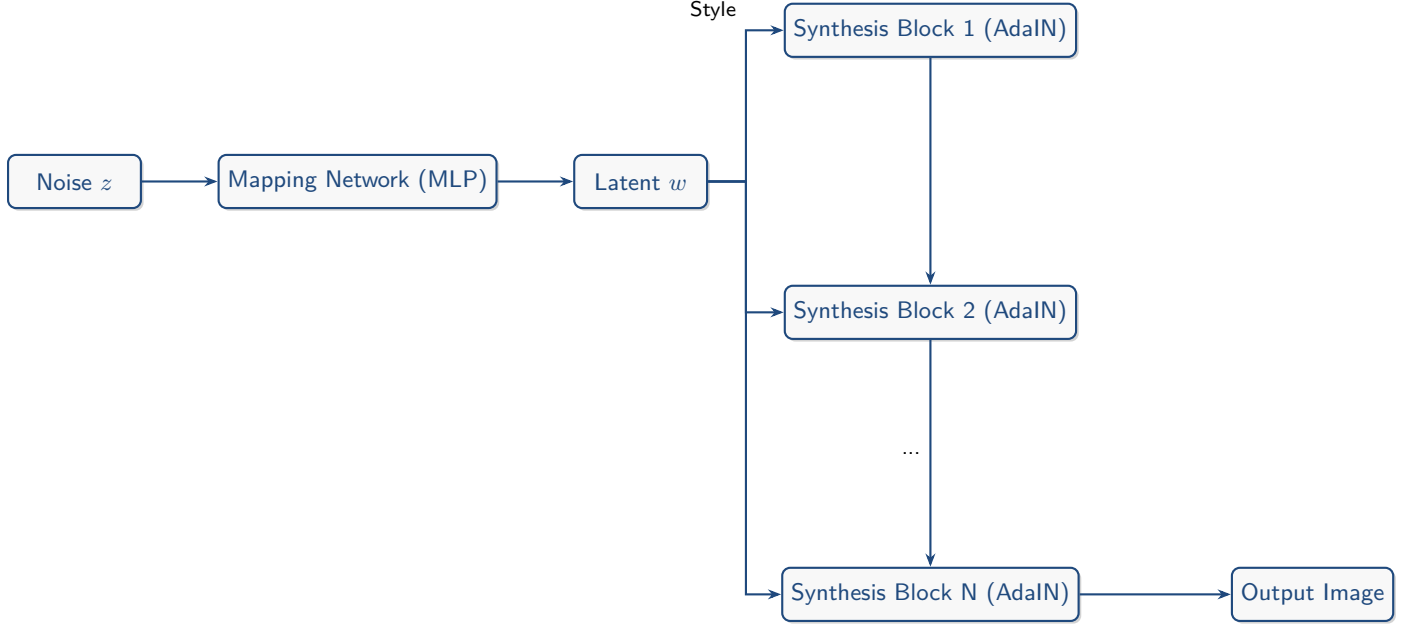


Figure 4: StyleGAN Architecture

- **Training Objective:** Standard GAN loss plus perceptual path length regularization to smooth latent space:

$$R = \mathbb{E}[\|J_w^T y\|_2], \quad \text{where } J_w \text{ is Jacobian} \quad (7)$$

- **Use Cases:** Photorealistic faces (FFHQ dataset), art generation, style transfer/mixing (e.g., blending facial attributes).

5.2 CycleGAN

CycleGAN (Zhu et al., 2017) enables unpaired image-to-image translation.

- **Architecture:** Two generators ($G : X \rightarrow Y$, $F : Y \rightarrow X$) and discriminators (D_X, D_Y). Cycle consistency ensures translations are invertible.

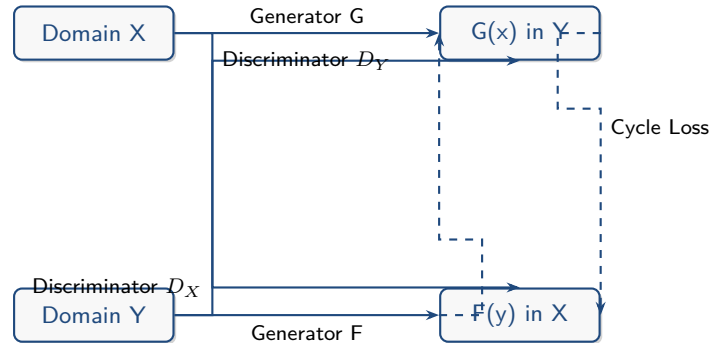


Figure 5: CycleGAN Architecture

- **Training Objective:** Adversarial losses for G, D_Y and F, D_X , plus cycle consistency:

$$\mathcal{L}_{cyc} = \lambda(\|F(G(x)) - x\|_1 + \|G(F(y)) - y\|_1) \quad (8)$$

- **Use Cases:** Style transfer (photos to Monet paintings), object transfiguration (horses to zebras), medical imaging domain adaptation.

5.3 BigGAN

BigGAN (Brock et al., 2018) scales GANs for high-fidelity, diverse image generation.

- **Architecture:** Class-conditional (like CGAN), using large ResNet-based networks, shared class embeddings, and truncation trick (sampling truncated noise for quality-diversity trade-off).
- **Diagram:** Extends CGAN with deeper layers and self-attention.
- **Training Objective:** Hinge loss variant of conditional GAN:

$$\mathcal{L}_D = -\mathbb{E}[\min(0, -1 + D(x, y))] - \mathbb{E}[\min(0, -1 - D(G(z, y), y))] \quad (9)$$

Plus orthogonal regularization on weights to prevent explosion.

- **Use Cases:** Large-scale datasets like ImageNet at 512x512 resolution, conditional generation for diverse classes, benchmark for FID/IS metrics.

Model	Key Innovation	Objective Additions	Primary Use Cases
StyleGAN	Style-based generator	Path length reg.	Faces, style mixing
CycleGAN	Cycle consistency	L1 cycle loss	Unpaired translation
BigGAN	Scaling & truncation	Hinge loss, ortho reg.	High-res class-conditional

Table 2: Comparison of Advanced GANs

These models address Vanilla limitations: StyleGAN for control, CycleGAN for unpaired data, BigGAN for scale.

6 Conclusion

This report has elaborated on GAN fundamentals, challenges, variants, and advancements, grounded in the referenced repository. Future directions include hybrid models with diffusion and improved stability.